

Md Hussain

Senior AI Product Engineer | Full Stack Engineer – Voice AI, Agentic Systems, LLM Infrastructure

+91 8860767389 | hussainakhtar1111@gmail.com | hussnakhtr.com

linkedin.com/in/md-hussain-baa178136 | github.com/HUSS41N

ABOUT

Senior AI Product Engineer with nearly **5 years** of software engineering experience building production **LLM, voice AI** and **agentic** systems. Led the development of **YEAPP**, a voice-first social product for [KRAFTON India](#) (the **PUBG MOBILE** company), to the Google Play Store in under 40 days, built **Scribe**, an AI therapy-session scribe for **Narayana Health**, and built the **voice moderation system** for [Gameberry Labs](#)' **Ludo Star** with language support across all **Indic languages, Arabic** and **Mexican Spanish**.

Deep expertise in **LangChain, LangGraph, RAG, vector search, OpenSearch, Model Context Protocol (MCP), multi-agent orchestration, LLM evaluation** and **real-time voice AI (WebRTC, model routing, model benchmarking)**. Strong **full-stack** foundation in **TypeScript, React / Next.js, React Native, Node.js / NestJS** and **Python / FastAPI**, deployed on **Docker, Kubernetes** and **AWS**.

Owned the **scalable agentic AI platform** at Jivi AI (**200+ agents in production, 3.8M+ app downloads, 80–100k DAU**) with **real-time evaluation** reaching **96% accuracy on OSCE cases**, and built an **AI-assisted application security platform** at Sudoviz. Cut voice-call cost by **~75%** (INR 4–5 to INR 0.7–1.3 per minute) with a **cascading system** and infrastructure optimization.

TECHNICAL SKILLS

AI / LLM: **LangChain, LangGraph, MCP** (Model Context Protocol), Claude, Multi-Agent Systems, Agent Workflows, Tool Calling, **RAG, Vector Search**, Embeddings, Vector Databases, **OpenSearch**, Generative AI, Prompt Engineering, LLM Evaluation, Tracing & Observability

Voice AI: Real-Time Voice Agents, Voice Moderation, Speaker Diarization, Ambient Audio Processing, STT / TTS Pipelines, **WebRTC**, Speech Workflows, Model Routing, Model Benchmarking, Local & Hosted Model Inference

Languages & Backend: Python, TypeScript, JavaScript, C++, Node.js, NestJS, Express.js, FastAPI, Flask, Django, REST APIs, WebSockets, MongoDB

Frontend: React.js, Next.js, React Native, Micro-frontends, Design Systems

Infra & Core: Docker, Kubernetes, AWS, Azure AI, Microservices, System Design, Distributed Systems, Real-Time Systems, API & Database Design

EXPERIENCE

Senior AI Product Engineer – The Product Highway Apr 2026 – Present

Bengaluru, India

Tech: Python, TypeScript, Node.js, NestJS, React, React Native, LangChain, LangGraph, MCP, Claude, WebRTC, WebSockets, Speaker Diarization, Docker, AWS

- Led the development of **YEAPP**, a voice-first social media product for [KRAFTON India](#) (the **PUBG MOBILE** company), leading a 4-engineer team and owning the **voice AI infrastructure** and the **React Native** frontend and **NestJS** backend services; took it from development to live on the Google Play Store in under 40 days.
- Architected [Echo](#), a **voice AI platform** to build, deploy, benchmark, and orchestrate **voice agents** and workflows, with a unified model-access layer spanning **88+ vendors** and **440+ local and hosted models**.
- Built the **voice moderation system** for [Gameberry Labs](#)' **Ludo Star**, with language support across all **Indic languages, Arabic** and **Mexican Spanish**.
- Built **Scribe** for **Narayana Health**, an AI **mental-health therapy scribe** with **speaker diarization, room ambience listening**, automated clinical **note writing**, and secure preservation of session context and records.
- Built an **MCP server** for [Echo](#), the voice AI platform, so anyone can connect it to **Claude** and create, configure and deploy **voice agent flows** through natural language.
- Built a **model benchmarking and evaluation** framework measuring latency and cost per model, feeding an automated **model router**.
- Cut voice-calling cost for [ANSR Global](#) from INR 4–5/min to **INR 0.7 -- 1.3/min (~75%)** with a **cascading system** that adapts to each call's requirements, plus **infrastructure optimization**.
- Delivered production **voice agents** for **healthcare AI intake** and **AI interviewing**, running streaming speech-to-text, LLM reasoning and text-to-speech pipelines over **WebRTC**.
- Authored an internal **frontend design skill** and design-system guidelines for consistent, maintainable AI-generated frontend code.

Software Development Engineer II – Jivi AI May 2024 – Apr 2026

Gurugram, India

Tech: Python, LangGraph, LangChain, MCP, FastAPI, OpenSearch, React, Next.js, MediaPipe, Docker, Kubernetes, AWS

- Owned the agentic AI systems behind Jivi's healthcare app (**3.8M + downloads, 80–100k DAU**), architecting a scalable platform that runs **200+ agents in production**.
- Built a **node-based multi-agent workflow system** with agent orchestration, **tool calling** with **MCP** tool integrations, conditional flows, and auto-generated REST APIs per workflow, using **LangGraph** and **LangChain**.
- Built **LLM evaluation** for every agentic system: large-scale pipelines for bulk workflow testing, response-quality scoring and prompt evaluation, plus **real-time evaluation** that combines **agentic eval rubrics** with signals from real clinical systems; agents reached **96% accuracy on OSCE cases**.
- Designed **LLM memory architectures** (short-term, long-term, and episodic) to improve context retention and multi-turn reasoning.
- Built **agent versioning** and controlled rollout mechanisms with **A/B testing** for safe production deployments.
- Built **semantic product search** and recommendation systems using **embeddings, vector search** and **OpenSearch** across healthcare datasets.
- Built a real-time, camera-based **heart-rate monitoring** system with **MediaPipe**, estimating heart rate from facial and finger signals with **99.2% accuracy**, benchmarked against a real **pulse oximeter**.
- Architected **micro-frontends** and a unified **design system** focused on performance, modularity, and maintainability.

Senior Software Engineer – Sudoviz Mar 2023 – May 2024

California, United States (Remote)

Tech: TypeScript, Node.js, Python, Azure AI, OAuth, GitHub API, VS Code Extension API

- Owned engineering for Sudoviz, an **AI-assisted application security platform**, building the backend services, APIs, integrations and developer tooling end to end.
- Built **LLM-powered code-analysis workflows** on **Azure AI** that scan GitHub repositories, detect and triage security vulnerabilities, and generate remediation guidance.
- Built a **VS Code extension** that surfaces vulnerabilities inline and lets developers investigate and remediate issues without leaving the IDE.
- Built **AST-based code analysis** for vulnerability detection and streamed **zero-knowledge audit events** to SIEMs (Splunk, Datadog).
- Built the enterprise **control plane** dashboard: vulnerability inventory, security graph, telemetry volume and risk scores.
- Optimized backend APIs by eliminating redundant requests; implemented **role-based access control** and **OAuth** integrations for secure repository connectivity.

Software Engineer – PakkaProfile Jan 2022 – Nov 2022

Bengaluru, India

Tech: React, React Native, Next.js, Node.js, REST APIs, Microservices, Data Processing Pipelines, MongoDB, Unity

- Shipped **15+ psychometric games** with **React** and **React Native** for a game-based hiring and assessment platform, at a scale of **150k+ users**.
- Built reusable game and assessment UI components shared across **React** and **React Native**, keeping behavior consistent across web and mobile and speeding up delivery of new games.
- Designed backend **microservices** and **data-processing pipelines** that ingest gameplay data (choices, timing, responses) and compute personality, cognitive and language assessment scores and reports, served to candidate-facing flows via **REST APIs**.
- Implemented a **real-time navigation system** (mini-map, compass and point-to-point guidance) in an interactive **Unity** environment.
- Delivered a **100% performance improvement** on the **Next.js** website, roughly 2x faster page loads.

EDUCATION

Master of Computer Applications (MCA) 2020 – 2022

Jamia Hamdard, New Delhi, India

Relevant coursework: Data Structures & Algorithms, Operating Systems, Database Management Systems, Computer Networks, Software Engineering, Artificial Intelligence

Bachelor of Computer Applications (BCA) 2017 – 2020

Jamia Hamdard, New Delhi, India

Relevant coursework: Programming (C, C++, Java), Data Structures, Database Systems, Web Technologies, Computer Organization